



Comparing Optimization Algorithms in the Fitting of Linear Mixed Models: Evaluating Speed and Accuracy using lme4 in R and lmm in Julia

Douglas Bates

University of Wisconsin-Madison

Colin Longhurst

University of Wisconsin-Madison

Abstract

The **Timings** package allows for the comparison of several optimizers, in both R and Julia, used in the fitting of various linear mixed models. In R the optimizers are called by lmer from the **lme4** package (version 1.1-8). In Julia the optimizers are called by lmm from **MixedModels** package. From the **Timings** package, conclusions regarding an optimizers relative speed, accuracy and general effectiveness of different optimizers paired with different types of models (ranging from simple to complex) can easily be drawn and interpreted.

There are differences in the model formulations in **lme4** and in **MixedModels**. The numerical representation of the model in **lme4** and the method of evaluating the optimizers, described in this paper, is the same for all models. In **MixedModels** there are specialized representations for some model forms, such as models with a single grouping factor for the random effects. Some of the specialized representations allow for evaluation of the gradient of the objects, which can enhance convergence (but, interestingly, sometimes can impede convergence).

Keywords: optimizers, mixed models, linear mixed models, lme4, lmm, R, Julia.

1. Introduction

Mauris ac lectus sagittis, mollis diam porttitor, efficitur purus. Duis ut eros vitae dolor vulputate sagittis. Aenean viverra, dui in interdum consectetur, lectus elit placerat urna, dictum ultricies ex justo eget est. In at leo quam. Suspendisse eros nulla, gravida nec suscipit vel, egestas eget neque. Etiam pharetra, nisl nec viverra ultricies, mauris lectus facilisis velit, vel aliquam ipsum lorem ut sem. Nulla ornare, justo at tempus blandit, nulla urna porta ligula, id dictum mi est ac massa. Quisque at lacus neque. Pellentesque est nulla, dignissim

eget lobortis ut, vehicula at erat. Aenean ornare lacus mattis, elementum elit vel, tempor risus. In in purus tempor lacus imperdiet rhoncus nec nec tortor. Duis sagittis nisl ante, id egestas neque tristique fermentum. Fusce aliquet, odio non auctor aliquet, purus orci venenatis purus, sit amet pulvinar nisi est at dolor. Pellentesque lobortis dui eros, et ultricies tellus ultrices at. Ut sit amet interdum justo. Integer placerat vehicula interdum.

2. Methods

To provide consistency we have copied all the data sets used in the timings to the **Timings** package itself. We have done all timings on the same computer. This computer has a relatively recent Intel processor and we used the **Intel Math Kernel Library (MKL)** with **Julia**. We attempted to use **Revolution R Open (RRO)** as the R implementation as it can be configured with **MKL**. However, we ran into version problems with this so we used the standard Ubuntu version of R linked against OpenBLAS, which is also multi-threaded.

Variables were renamed in the pattern:

- **Y** the response
- **A, B, ...** categorical covariates
- **G,H, I, ...** grouping factors for random effects
- **U, V, ...** (skipping **Y**) continuous covariates

The timing results are saved in **JSON (JavaScript Object Notation)** files in the directory accessible as

```
> system.file("JSON",package="Timings")
```

within **R**. The directory name will end with `./Timings/inst/JSON/` in the package source directory, for example the result of cloning the [github repository](#). There is one `.json` file for each data set. Each such file contains results on timings of one or more models.

The **Timings** package for R provides a `retime` function that takes the name of one of these JSON files and, optionally, the name of a file with the updated timings. Similarly there are some source files for **Julia** retimings.

```
> include("../julia/retime.jl")
> retime("../JSON/Alfalfa.json", "/tmp/Alfalfa.json")
> retime("../JSON/Alfalfa.json", "/tmp/Alfalfa.json")
```

INCLUDE SUMMARY TABLE USING `res.rda`

The timing was repeated so that compilation time is not included in the results. This repetition is only needed once per session.

A careful examination of these results shows that the main differences in the **Julia** timings (the **R** timings are merely reported, not evaluated) are that the `LN_BOBYQA` and `LD_MMA` optimizers

are much faster in the second run. This is because much of the code needs to be compiled the first time that a derivative-free optimizer and a derivative-based optimizer are used.

The names of the optimizers used with `lmm` are those from the **NLopt** package for **Julia**. Names that begin with `LD_` are gradient-based methods. Names that begin with `LN_` are derivative-free methods. There is one other derivative-free method, `LN_PRAXIS`, available in the **NLopt** package but, for some reason, it can hang on very simple problems like this. Frequently we omit it.

The optimizers used with `lmer` include the `Nelder_Mead` optimizer built into the **lme4** package, the `bobyqa` optimizer from the **minqa** package, the derivative-free optimizers from the **nloptr** package and several optimizers from the **optimx** package.

The `optimx:bobyqa` optimizer is just a wrapper around `bobyqa` (bounded optimization by quadratic approximation) from the **minqa** package and should provide results similar to those from the `bobyqa` optimizer. For some reason the number of function evaluations is not reported for the version in **optimx**.

The optimizers from **nloptr** (i.e. those whose names begin with `NLOPT_LN_`) use the same underlying code as do the similarly named optimizers in the **NLopt** package for **Julia**. The number of iterations to convergence should be similar for the same underlying code, although not necessarily exactly the same because the evaluation of the objective in **R** and in **Julia** may produce slightly different answers. Also the convergence criteria in the **Julia** version are more strict than those in the **R** version.

Also shown are the value of the criterion (negative twice the log-likelihood, lower is better) achieved, the elapsed time and the number of function and gradient evaluations. The `nopt` value is the number of parameters in the optimization problem. `mtype` is the model type in the **Julia** code. There are special methods for solving the penalized least squares (PLS) problem, and for evaluating the objective and its gradient when there is only one grouping factor for the random effects. The model type is called `PLS0ne`.

The **Alfalpa** example is a particularly easy one and all of the optimizers converge to an objective value close to -10.81023 in less than 0.6 seconds.

3. Results

For the **Alfalpa** data there is not much of a burden in refitting the model with all the **Julia** optimizers just to get the table shown above. But other examples can take an hour or more to converge and we don't really need to refit them every time. The `tabulate.jl` file contains a function `optdir` to create a **DataFrame** from the results of all the model fits.

```
> include("../julia/tabulate.jl")
> res = optdir("../JSON");
>
```

Row	opt	dsname	n	np	excess	time	reltime
-----	-----	-----	-----	-----	-----	-----	-----
1	"LD_CCSAQ"	"Alfalpa"	72	1	0.0	0.0017	27.5171
2	"LD_CCSAQ"	"AvgDailyGain"	32	1	0.0	0.0014	13.9236
3	"LD_CCSAQ"	"AvgDailyGain"	32	1	0.0	0.0014	10.8009

4	"LD_CCSAQ"	"BIB"	24	1	0.0	0.0013	11.7005
5	"LD_CCSAQ"	"Bond"	21	1	0.0	0.0009	13.6969
6	"LD_CCSAQ"	"bs10"	1104	20	0.0	1.0958	12.608
7	"LD_CCSAQ"	"bs10"	1104	8	39.9948	0.0375	23.2699
8	"LD_CCSAQ"	"cake"	270	1	0.0	0.0033	18.8753
9	"LD_CCSAQ"	"Cultivation"	24	1	0.0	0.0009	15.6477
10	"LD_CCSAQ"	"Demand"	77	2	3.21928	0.0055	22.0827
11	"LD_CCSAQ"	"dialectNL"	225866	6	0.0	6.9896	3.942
.							
.							
.							
19	"LD_CCSAQ"	"gb12"	512	8	103.176	0.0218	13.1358
20	"LD_CCSAQ"	"HR"	120	3	0.0	0.0089	8.6481
21	"LD_CCSAQ"	"Hsb82"	7185	1	192.73	0.0102	3.7438
22	"LD_CCSAQ"	"IncBlk"	24	1	0.55726	0.001	5.0422
23	"LD_CCSAQ"	"kb07"	1790	72	8.20739	17.4698	9.2732
24	"LD_CCSAQ"	"Mississippi"	37	1	0.93471	0.0006	15.9582
25	"LD_CCSAQ"	"mm0"	69588	6	0.0	4.8286	1.1034
26	"LD_CCSAQ"	"Oxboys"	234	3	136.788	0.0169	5.3039
27	"LD_CCSAQ"	"PBIB"	60	1	0.0	0.0014	8.6177
28	"LD_CCSAQ"	"Penicillin"	144	2	0.0	0.0131	15.3763
29	"LD_CCSAQ"	"Semiconductor"	48	1	0.0	0.0012	1.133
30	"LD_CCSAQ"	"SIMS"	3691	3	3.60856	0.134	2.3154

3.1. Proportion Converged

Etiam placerat commodo dapibus. Vivamus a porta dui. Phasellus accumsan sollicitudin cursus. In nec lorem sit amet purus vulputate suscipit quis ac lectus. Aliquam sollicitudin malesuada blandit. Aliquam sed nisi accumsan, pulvinar diam eget, facilisis urna. Fusce a dignissim ipsum. Integer maximus, felis nec lobortis vulputate, nunc leo blandit sem, id suscipit velit leo in purus. Pellentesque vel arcu sit amet nulla tincidunt ornare. Mauris sodales lectus eget est laoreet lobortis.

3.2. Reliability

Sed iaculis sodales elit quis vehicula. In et tristique neque, sodales aliquet metus. In posuere dictum nisl, quis laoreet augue congue a. Aenean in commodo neque, sit amet hendrerit ex. Aliquam id faucibus ante. Vivamus in fermentum nunc. Nam condimentum eros id orci pretium, quis aliquam magna eleifend.

4. Conclusions

5. References

Note all optimization packages used, lmm, lme4, jsonlite,

Affiliation:

Douglas Bates
Emeritus Professor
Department of Statistics
University of Wisconsin-Madison
1300 University Ave Madison, WI 53706-1685 U.S.A.
E-mail: dmbates@stat.wisc.edu
Github URL: <https://github.com/dmbates>